

# 1. Introduction

---

## 1.1 Purpose

This document provides the architectural and design details of the TelcoPulse real-time analytics system for monitoring mobile network metrics. The system processes real-time streams from telecom providers, derives insights, and visualizes them through a live dashboard.

## 1.2 Scope

The solution focuses on real-time ingestion, transformation, storage, and visualization of network performance data using AWS-native services. It includes:

- Real-time data ingestion via Amazon Kinesis
- Data transformation with AWS Glue and Spark Streaming
- Storage in Amazon S3
- Querying via AWS Athena
- Visualization using Streamlit on Amazon ECS

## 1.3 Definitions, Acronyms, and Abbreviations

| Term | Definition                |
|------|---------------------------|
| KPI  | Key Performance Indicator |
| ECS  | Elastic Container Service |
| S3   | Simple Storage Service    |
| GPS  | Global Positioning System |
| ETL  | Extract, Transform, Load  |
| SQL  | Structured Query Language |

## 1.4 References

- AWS Documentation
- Apache Spark Documentation
- Streamlit Documentation
- Project README

## 1.5 Overview

The document outlines a modular architecture using AWS-native services to ensure scalability, reliability, and low-latency analytics on mobile network metrics.

## 2. Architectural Goals and Constraints

---

### 2.1 System-wide Design Decisions

- Real-time streaming architecture
- Serverless and fully managed services where possible
- Decoupling components for maintainability and scalability

### 2.2 Architectural Constraints

- AWS-only implementation
- Near real-time latency (<5 minutes)
- Autoscaling components to support variable input loads

## 3. Use-Case View

---

### 3.1 Use-Case Model

Actors:

- Data Engineering Team
- Network Operations Analysts

Use Cases:

- Monitor signal strength trends
- Identify network outages by postal code
- Analyze GPS accuracy by operator

## 4. Logical View

---

## 4.1 Major Logical Components

- **Data Ingestor** (Amazon Kinesis)
- **Stream Processor** (AWS Glue with Spark)
- **Data Lake** (Amazon S3)
- **Data Catalog** (Glue Crawlers)
- **Query Engine** (AWS Athena)
- **Dashboard UI** (Streamlit on ECS)

## 5. Process View

---

### 5.1 Concurrent Processes

- Kinesis continuously streams data
- Spark Streaming processes events in micro-batches
- ECS runs containerized dashboard with auto-refresh logic

## 6. Implementation View

---

### 6.1 Development Environment

- Python 3.x for Spark and Streamlit
- AWS SDK / Boto3
- AWS CLI, Terraform

### 6.2 Configuration Management

- Git for source control
- Versioning of ETL jobs and dashboards via tags/branches

## 7. Data View

---

### 7.1 Storage & Retrieval

- Raw and processed data in S3

- Athena queries provide real-time metrics to dashboard

## 7.2 Data Integrity and Security

- S3 bucket policies for access control
- IAM roles for Glue, Athena, ECS
- Encryption at rest and in transit enabled

# 8. Size and Performance

---

## 8.1 Performance Benchmarks

- Data latency: < 5 minutes from ingestion to visualization
- Throughput: Scalable with stream size via Kinesis shard scaling

## 8.2 Scalability Constraints

- Number of Kinesis shards
- Glue job parallelism
- ECS service auto-scaling settings

# 9. Quality

---

| Attribute       | Strategy                               |
|-----------------|--|
| Usability       | Interactive, auto-refresh dashboard UI |
| Reliability     | Managed services with fault tolerance  |
| Performance     | Micro-batch processing and tuning      |
| Maintainability | Modular service design                 |
| Security        | IAM, encryption, private networking    |

# 10. Appendices

---

# 10.1 Revision History

| Date       | Version | Description             |
|------------|---------|-------------------------|
| 2025-05-15 | 1.0     | Initial draft generated |